

Same Goal, Different Words: Spatial Instruction Sensitivity in World–Action Models

Ali-Adeeb Abbas^{1,2}, Esmail Seraj², Behrad Toghi², and Taskin Padir^{1,3*}

Abstract: A robot that follows one spatial instruction may fail when the same goal is described from another object’s perspective. We examine this sensitivity in three world–action models through 864 rollouts in four RoboLab scenes. We keep the initial scene and intended manipulation fixed while changing “move the mustard so that it is right of the box” to “move the mustard so that the box is left of it.” Describing the relation through the reference object reduces the proportion of runs reaching a stable placement by 12.2 percentage points on average, with the effect concentrated in particular relations. Inspecting executions also shows that a spatial score can reward moving the wrong object, while final-state checks reveal that many attained arrangements are subsequently lost. These findings make instruction consistency a necessary complement to task performance: evaluation should test alternative descriptions of the same goal and verify the requested manipulation, not only the resulting spatial relation.

Keywords: World–Action Models, Spatial Language, Robot Evaluation

1 Introduction

World–action models (WAMs) generate robot actions together with predictions of future observations [1]. Their usefulness as general-purpose policies depends on turning a user’s description into the intended manipulation. For a spatial placement task, this requires identifying both the object to move and its destination. A successful rollout under one description leaves a basic question unanswered: does the policy respond to the requested goal consistently, or does success depend on how that goal is expressed?

Language-conditioned robot policies have been examined from several complementary perspectives. TRI’s LBM study tests language conditioning in visually similar scenes [2]. CAST connects diverse descriptions to atomic commands and uses counterfactual action supervision to improve instruction following [3]. LIBERO-CF tests responses to changed goals [4], while LIBERO-Plus, LIBERO-Para, and related work examine language perturbations, equivalent paraphrases, and irrelevant context [5, 6, 7]. Within WAMs, SG-WAM studies semantic guidance and language alignment [8]. These studies motivate examining whether a policy preserves a goal across descriptions, in addition to whether it responds when the goal changes.

We focus on a specific spatial transformation. Placing the mustard to the right of a box also places the box to the mustard’s left. Both instructions can explicitly require moving the mustard, so reversing the relation need not change the physical task. We test this equivalence in three released WAMs using repeated runs from identical initial scenes. The results show that successful placement does not imply consistency across equivalent descriptions, and that this sensitivity varies across relations. A second finding concerns evaluation itself: a spatial relation can become true without the requested object being moved, or cease to hold under continued control. Together, these findings support evaluating *which manipulation is preserved across descriptions*, rather than treating a single placement score as sufficient evidence of instruction following.

^{*1}RIVeR Laboratory, Northeastern University, Boston, MA, USA. ²General Motors Robotics, Mountain View, CA, USA. ³Amazon Robotics, North Reading, MA, USA. Correspondence: abbas.a@northeastern.edu.

Table 1: One physical goal described three ways. The robot must move the mustard bottle (the *target*) relative to the raisin box (the *reference*). Only the instruction changes between runs.



Goal: mustard to the right of the box.
All three instructions identify the mustard as the object to move. S and I differ in which object is the subject of the spatial clause.

Form	Exact instruction
Direct (D)	Place the mustard bottle on the table to the right of the raisin box.
Target-first (S)	Move the mustard bottle on the table so that the mustard bottle is to the right of the raisin box.
Reference-first (I)	Move the mustard bottle on the table so that the raisin box is to the left of the mustard bottle.

2 Paired Spatial Evaluation

To test whether wording changes behavior without changing the task, we repeat each placement from the same physical start under three instructions. We compare how often the requested arrangement is reached, then examine whether that score captures the intended manipulation and its final outcome.

We use four tabletop scenes from RoboLab, a simulation benchmark for generalist manipulation policies [9]. The cube scene asks the robot to place a cube left, right, in front of, or behind a bowl. The butter and mustard scenes each ask it to place that object left, right, or on top of a raisin box. The bowl scene asks it to stack either bowl on the other. Seven goals are provided by RoboLab, and five extend these same scenes without changing their objects. Appendix A shows the scenes and all twelve goals.

Table 1 introduces our instruction forms. D is a concise command, using the benchmark’s supplied instruction for its original goals. S expands that command while keeping the target as the subject of the spatial clause. I instead describes the arrangement through the reference object and the opposite relation, such as left instead of right. S versus I therefore tests reference reversal with the surrounding sentence held fixed. D versus S tests the expansion itself. D versus I combines both changes and is reported in Appendix C. For stacking, “underneath and supporting” does not explicitly require the nesting checked by the scorer. We keep these conditions visible and examine the result without them.

We evaluate the DROID versions of Cosmos3 Nano, Cosmos3 Edge, and FLUX 3 Action [10, 11, 12]. Nano and FLUX perform strongly on RoboLab’s leaderboard, while Edge offers a smaller Cosmos-family comparison [13]. Each receives wrist and two exterior camera images plus the robot’s proprioceptive state. We vary object positions and orientations on the tabletop and retain eight starts per scene that a scripted controller can solve. For every model and goal, the simulator is reset to exactly the same start before each instruction. These are our *paired runs*: the robot state, object poses, and initial observations agree, while the wording changes. Three models, twelve goals, three descriptions, and eight starts give 864 episodes. Each runs for 20–40 seconds, including continued control after placement. Appendix B gives the full settings.

We measure the proportion of runs that reach a *stable placement*. The requested relative arrangement must hold for at least one second while the target is supported, moves slower than 2 cm/s, and is not touching the gripper. Directions are defined in the robot’s coordinate frame. We also check the same conditions during the final second, distinguishing an arrangement reached earlier from one retained until the end. Neither check identifies which object was moved. A cube can end up behind a bowl because the robot moved the cube, or because it moved the bowl.

For each pair, we compare these binary outcomes and average first across goals within each scene, then equally across scenes and models. Thus, a positive S–I gap means more runs pass under S than under I. Uncertainty accounts for the shared physical starts rather than treating all rollouts as independent. Three goals already hold at the start. Figure 2 marks them, and we also report rates excluding them. Statistical details and all counts appear in Appendices B.1 and C.4.

3 Results

Task competence under one description does not guarantee consistency across equivalent instructions. Reference reversal reduces the stable-placement rate by 12.2 percentage points on average, with the same direction for all three models (Figure 1, 95% interval 8.6–15.4, adjusted $p < 0.001$). Because S and I share the expanded sentence structure, this comparison tests relational reformulation without also changing from a direct to an expanded command. With the physical task unchanged, this result exposes sensitivity to its description, although argument order and relation vocabulary are not separated.




Figure 1: Wording effects across 32 physical starts. Dots show each start’s average outcome difference. Diamonds and bars show means and 95% intervals. Positive S–I favors target-first wording. Wider sections indicate more starts with similar effects.

The effect depends on the object–relation combination. Mustard-right succeeds in 20 of 24 target-first runs but only two reference-first runs (Figure 3). In contrast, the cube scene’s zero average masks effects with different signs across directions, including an advantage for I on cube-behind (Figure 2). Aggregate scores can therefore conceal inconsistent responses to individual goals. The positive pooled gap persists without bowl stacking, whose descriptions have a nesting ambiguity (Appendix B.1).




Figure 2: Stable-placement rates (%). Overall columns average goals within scenes and weight scenes equally. Each row contains 96 episodes. Individual goals use eight starts. Hatched goals already hold initially and remain included in the overall rates. These scores check the arrangement, not which object moved.

Edge · mustard right of the raisin box



Figure 3: Edge succeeds under target-first (S) wording but fails under reference-first (I), despite the same initial scene and requested mustard placement. This pair was selected for its contrasting outcomes. Additional views appear in Appendix D.

Reaching a scored relation is weaker than completing the intended manipulation. The overall placement rate falls from 43.2% during execution to 21.8% at episode end. Excluding initially satisfied goals gives 26.8% and 15.3%. In raw counts, 192 of 383 runs that attain a stable placement lose it before the end. Separately, a Nano cube-behind rollout passes by moving the bowl rather than the requested cube (Appendix D). Spatial success can therefore overstate instruction following. Evaluation must distinguish the arrangement reached, the object manipulated, and persistence under the chosen stopping rule.

4 Discussion

Instruction consistency should be measured alongside task performance. A policy can perform a placement under one description yet fail to recover the same goal from another. Following SIMPLER’s controlled-simulation approach [14], our comparison keeps the physical start fixed and reveals this distinction directly. Testing several descriptions per goal would reduce dependence on favorable phrasing and identify failures to target during training.

The relation-specific pattern cautions against a universal prompting rule. Target-first wording performs better on average, but opposing effects within the cube scene show why the average cannot guide every task. One possible explanation is that policies rely on familiar object–relation combinations more than their logical equivalence. Separately varying vocabulary and argument order would test this interpretation. For WAMs, generating a future observation does not itself demonstrate reliable goal interpretation. Appendix E reports why our forecast analysis did not resolve the source of the behavioral failures.

Evaluation must also preserve the requested action’s semantics. As with LIBERO-CF’s separation of object contact and task success [4], a geometric relation captures only part of instruction following. Checking which object moved and whether the arrangement remains at the end addresses two concrete gaps exposed here. Broader scenes and descriptions are needed to establish their prevalence, and this study does not compare WAMs against action-only policies.

5 Conclusion

Across three WAMs, equivalent spatial descriptions produce inconsistent placement outcomes, while spatial scores can conceal wrong-object manipulation and arrangements lost during execution. Evaluating several descriptions per goal, together with object identity and the final arrangement, exposes a gap between achieving a placement and reliably following the instruction.

References

- [1] M. J. Kim, Y. Gao, T.-Y. Lin, Y.-C. Lin, Y. Ge, G. Lam, P. Liang, S. Song, M.-Y. Liu, C. Finn, and J. Gu. Cosmos policy: Fine-tuning video models for visuomotor control and planning. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=wPEIStHxYH>.
- [2] TRI LBM Team. A careful examination of large behavior models for multitask dexterous manipulation. *arXiv preprint arXiv:2507.05331*, 2025.
- [3] C. Glossop, W. Chen, A. Bhorkar, D. Shah, and S. Levine. CAST: Counterfactual labels improve instruction following in vision-language-action models. *arXiv preprint arXiv:2508.13446*, 2025.
- [4] Y. Fang, Y. Feng, D. Jing, J. Liu, Y. Yang, Z. Wei, D. Szafrir, and M. Ding. When vision overrides language: Evaluating and mitigating counterfactual failures in VLAs. *arXiv preprint arXiv:2602.17659*, 2026.
- [5] S. Fei, S. Wang, J. Shi, Z. Dai, J. Cai, P. Qian, L. Ji, X. He, S. Zhang, Z. Fei, J. Fu, J. Gong, and X. Qiu. LIBERO-Plus: A progressive robustness benchmark for visual-language-action models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 38574–38583, 2026.
- [6] C. Kim, M. Kim, M. Kang, H. Kim, and D. Jung. LIBERO-Para: A diagnostic benchmark and metrics for paraphrase robustness in VLA models. *arXiv preprint arXiv:2603.28301*, 2026. URL <https://arxiv.org/abs/2603.28301>. Accepted to EMNLP 2026, main conference.
- [7] A. Moskalenko, D. Pugacheva, D. Shepelev, A. Kuznetsov, V. Shakhuro, and E. Tutubalina. Bring the apple, not the sofa: Impact of irrelevant context in embodied AI commands on VLA models. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 4: Student Research Workshop)*, pages 840–860. Association for Computational Linguistics, 2026. doi:10.18653/v1/2026.eacl-srw.63. URL <https://aclanthology.org/2026.eacl-srw.63/>.
- [8] J. He, J. Li, Z. Zhong, H. Yan, R. Li, Y. Zheng, J. Zhu, T. Zhang, Y. Du, W. Chen, S. Zhou, and H. Li. SG-WAM: Text-grounded and spatial-aware semantic guidance for world-action models. *arXiv preprint arXiv:2608.08839*, 2026. URL <https://arxiv.org/abs/2608.08839>.
- [9] X. Yang, R. Dagli, A. Zook, H. Hadfield, A. Goyal, S. Birchfield, F. Ramos, and J. Tremblay. RoboLab: A high-fidelity simulation benchmark for analysis of task generalist policies. In *Proceedings of Robotics: Science and Systems*, 2026. doi:10.15607/RSS.2026.XXII.096.
- [10] NVIDIA. Cosmos3-Nano-Policy-DROID. Hugging Face model card and checkpoint, 2026. URL <https://huggingface.co/nvidia/Cosmos3-Nano-Policy-DROID>. Accessed September 27, 2026.
- [11] NVIDIA. Cosmos3-Edge-Policy-DROID. Hugging Face model card and checkpoint, 2026. URL <https://huggingface.co/nvidia/Cosmos3-Edge-Policy-DROID>. Accessed September 27, 2026.
- [12] Black Forest Labs. FLUX 3 Action DROID. Hugging Face model card and checkpoint, 2026. URL <https://huggingface.co/black-forest-labs/flux-3-action-droid>. Accessed September 27, 2026.
- [13] RoboLab Team. RoboLab-120 Leaderboard, 2026. URL <https://research.nvidia.com/labs/srl/projects/robolab/leaderboard.html>. Accessed September 27, 2026.
- [14] X. Li, K. Hsu, J. Gu, O. Mees, K. Pertsch, H. R. Walke, C. Fu, I. Lunawat, I. Sieh, S. Kirmani, S. Levine, J. Wu, C. Finn, H. Su, Q. Vuong, and T. Xiao. Evaluating real-world robot manipulation policies in simulation. In *Proceedings of the 8th Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, pages 3705–3728. PMLR, 2025.

A Evaluation details and full instructions

The comparison changes the instruction while fixing the requested placement and the object the robot should move. For each physical starting state, the same model is run separately under the direct (D), target-first (S), and reference-first (I) instructions. “Target-first” describes the order of objects in the relational clause after “so that”. Every instruction still names the same target object as the object to move. The control sequence is not fixed: each policy generates a new sequence from its instruction, so later observations may differ.

A.1 Scenes and starting states

The study includes eight starting states in each of four native RoboLab scenes. Small changes to object positions and orientations produce different starts, and a scripted controller checks the requested placements before learned-policy evaluation. The cube scene has four goals, each box scene has three, and the two-bowl scene has two, giving twelve goals and 36 instructions. With three models and eight starts for every goal, this yields 864 episodes. The planned two-bin scene did not produce a qualifying start with the scripted controller and was excluded before the learned-policy outcomes were collected. It contributes no observations to these results.

Table 2: Scenes, goals, and episode duration. All goals in a scene share its eight starting states. The already-satisfied column describes all eight starts in that scene.

Scene	Requested placements	Already satisfied	Duration
Cube and bowl	Cube left, right, in front, behind bowl	Cube right	30 s
Butter and raisin box	Butter left, right, on top of box	Butter right	40 s
Mustard and raisin box	Mustard left, right, on top of box	Mustard left	40 s
Two bowls	Either bowl on the other	None	20 s

Directions are defined in the robot’s coordinate frame. The left and right bowl identities are assigned from their positions at reset and remain fixed for the episode. Seven goals are supplied by RoboLab, and five add left/right placements using these same scenes and objects. In particular, the cube-right instruction is tested in the simple cube/bowl/banana scene used for the other cube goals, rather than in RoboLab’s separate cube-right scene with additional objects.




Figure 4: The four scene assets used in the study. These pilot views illustrate the objects and workspace. They are not the exact confirmation resets. All wording comparisons within a start use identical physical states.

A.2 Exact instructions

Tables below reproduce all tested instructions. D retains the benchmark task’s original wording where that goal exists on the scene. It is a baseline instruction, not evidence that wording came from a diverse human annotation process. S describes the relation starting with the target object, whereas I starts with the reference object and reverses the relation. The source file preserves a trailing space in the native mustard-on-box instruction. That whitespace is omitted from this typeset table.

A.2.1 Cube and bowl

Form	Instruction
Cube left of bowl	
D	Put the rubiks cube to the left of the bowl
S	Place the Rubik’s cube so that the Rubik’s cube is to the left of the bowl.
I	Place the Rubik’s cube so that the bowl is to the right of the Rubik’s cube.
Cube right of bowl	
D	Put the rubiks cube to the right of the bowl
S	Place the Rubik’s cube so that the Rubik’s cube is to the right of the bowl.
I	Place the Rubik’s cube so that the bowl is to the left of the Rubik’s cube.
Cube in front of bowl	
D	Put the rubiks cube in front of the bowl
S	Place the Rubik’s cube so that the Rubik’s cube is in front of the bowl.
I	Place the Rubik’s cube so that the bowl is behind the Rubik’s cube.
Cube behind bowl	
D	Put the rubiks cube behind the bowl
S	Place the Rubik’s cube so that the Rubik’s cube is behind the bowl.
I	Place the Rubik’s cube so that the bowl is in front of the Rubik’s cube.

Form	Instruction

A.2.2 Butter and raisin box

Form	Instruction
Butter box on raisin box	
D	Pick up the butter box and place it on top of the raisin box
S	Place the butter box so that the butter box is on top of and supported by the raisin box.
I	Place the butter box so that the raisin box is underneath and supporting the butter box.
Butter box left of raisin box	
D	Place the butter box on the table to the left of the raisin box.
S	Move the butter box on the table so that the butter box is to the left of the raisin box.
I	Move the butter box on the table so that the raisin box is to the right of the butter box.
Butter box right of raisin box	
D	Place the butter box on the table to the right of the raisin box.
S	Move the butter box on the table so that the butter box is to the right of the raisin box.
I	Move the butter box on the table so that the raisin box is to the left of the butter box.

A.2.3 Mustard and raisin box

Form	Instruction
Mustard bottle on raisin box	
D	Place the mustard on the raisin box.
S	Place the mustard bottle so that the mustard bottle is on top of and supported by the raisin box.
I	Place the mustard bottle so that the raisin box is underneath and supporting the mustard bottle.
Mustard bottle left of raisin box	
D	Place the mustard bottle on the table to the left of the raisin box.
S	Move the mustard bottle on the table so that the mustard bottle is to the left of the raisin box.
I	Move the mustard bottle on the table so that the raisin box is to the right of the mustard bottle.
Mustard bottle right of raisin box	
D	Place the mustard bottle on the table to the right of the raisin box.
S	Move the mustard bottle on the table so that the mustard bottle is to the right of the raisin box.
I	Move the mustard bottle on the table so that the raisin box is to the left of the mustard bottle.

A.2.4 Two bowls

Form	Instruction
Left bowl onto right bowl	
D	Stack the left bowl on the right bowl
S	Move the left bowl so that the left bowl is stacked on the right bowl.
I	Move the left bowl so that the right bowl is underneath and supporting the left bowl.
Right bowl onto left bowl	
D	Stack the right bowl on the left bowl
S	Move the right bowl so that the right bowl is stacked on the left bowl.
I	Move the right bowl so that the left bowl is underneath and supporting the right bowl.

The bowl task deserves separate interpretation. Its D instruction requests stacking, and its S instruction likewise specifies that the target is stacked on the other bowl. I describes the lower bowl as underneath and supporting the target bowl. That wording does not explicitly require the nesting checked by the benchmark predicate. The bowl results are therefore reported separately and the pooled wording contrast is also examined without them.

B Policy settings and evaluation criteria

The policies receive one camera view attached to the wrist, two fixed exterior views overlooking the workspace, and robot state. We retain RoboLab’s native camera configuration across all wording conditions. No camera is repositioned for a particular instruction. Each request returns 32 actions, executed at 15 Hz before the next observation and request. The final chunk is truncated at the episode horizon. The same request seed (6100) is used for the paired instructions. Episodes run for the complete duration in Table 2, even after a placement first passes the check. There are no automatic request retries. The policy implementation and camera preprocessing remain specific to each released checkpoint.

Table 7: Recorded inference settings. Values come from the policy server receipts used in the study.

Setting	Cosmos3 Nano	Cosmos3 Edge	FLUX 3 Action
Arithmetic	bfloat16	bfloat16	bfloat16
Sampling steps	4	4	4
Sampler	UniPC	UniPC	Cosmos UniPC
Noise shift	5.0	5.0	5.0
Visual guidance	3.0	3.0	4.0
Action guidance	—	—	1.0
Observation history	1	1	1
Actions per request	32	32	32
Control frequency	15 Hz	15 Hz	15 Hz
Prompt JSON formatting	No	Yes	No
Input/canvas setting	540 × 640	540 × 640	544 × 736

The Cosmos models use joint-position actions and a 480 resolution setting. Edge records a guidance interval of [960,1001]. FLUX uses its DROID camera layout and absolute action parameterization with the released action scale of 2.0. FLUX’s immutable configuration retains an inference-seed value of 0, while each inference call explicitly supplies the study seed 6100. This overrides the configuration value. Canvas dimensions are model-specific preprocessing settings, not different simulated camera placements. Full recorded configurations and release identities are preserved in the accompanying reproducibility files. The compact FLUX receipt does not contain a revision identifier, so a complete executed-weight manifest remains unavailable.

B.1 Placement criteria and statistical comparison

For left/right/front/behind, the spatial predicate uses the native 45-degree relation cone in the robot frame and requires the target to be supported by the table. For on-top placement, the native support predicate checks upward support and the target centroid within the reference footprint, with a 1 cm tolerance. Bowl stacking combines the native open-top containment predicate with contact between the bowls. In every case the gripper must be detached from the target. A stable placement requires the full predicate and target speed below 2 cm/s for one continuous second. Speed is computed by a backward difference of successive recorded positions at 15 Hz, which avoids treating solver velocity jitter at an unchanged position as motion.

The any-time outcome asks whether this condition holds anywhere in the episode. The episode-end outcome requires it throughout the final second. Neither condition proves that the intended object was manipulated. For example, moving the reference can make the requested relative arrangement true while leaving the target untouched. They are placement outcomes rather than a complete instruction-following score.

The saved initial states determine whether a requested arrangement holds before any policy action. Cube-right, butter-right, and mustard-left are initially satisfied in all eight starts. Those three goals account for 216 episodes, and the other nine account for 648 episodes. Every initially-satisfied episode passes the any-time criterion, so including these episodes increases absolute any-time rates without changing their S-I contribution, which is zero. Earlier per-episode labels disagreed with the saved-state classification for 44 episodes. The supplementary table uses the common saved-state labels consistently across forms and models. Original any-time and episode-end outcomes are unchanged.

For a wording comparison, the two binary placement outcomes are subtracted for the same model, starting state, and goal. For example, if S reaches the placement and I does not, the paired difference is +1. The reverse gives -1, and agreement gives 0. Differences are averaged over goals within each scene, then equally over scenes and models. This prevents scenes with more goals from dominating the pooled result. Confidence intervals use 10,000 bootstrap draws of the eight physical starts within each scene. All associated goals, models, and instruction forms remain together. The original pooled S-I and D-S tests use paired sign flips and Holm correction across those two tests. Per-model, subgroup, per-relation, terminal-outcome, and newly added D-I summaries are descriptive. They do not expand that original test family.

The pooled S-I estimate is 12.15 percentage points for any-time placement and 11.20 points for episode-end placement. Removing the bowl scene and giving each remaining scene equal weight leaves an any-time estimate of 10.65 points. This descriptive sensitivity check shows that the main direction does not depend on the bowl-language ambiguity. It does not establish equivalence of the bowl prompts. D-S is 4.08 points, with a 95% interval of $[-0.2, 8.4]$ and an adjusted $p = 0.089$. Failure to detect that difference is not evidence that instruction expansion has no effect.

C Additional behavioral results

C.1 Direct versus reference-first wording

The direct baseline can also be compared with reference-first wording. This added comparison is useful for connecting the controlled S-I intervention to the short instruction normally supplied with a task, but D and I also differ in length and sentence construction. It therefore does not isolate reference reversal as closely as S versus I. Table 8 and Figure 5 report it descriptively.

C.2 Relation-specific effects

Figure 6 shows that the average wording effect does not represent a uniform drop across relations. Mustard-right exhibits the largest contrast: 20 of 24 target-first episodes reach the placement, compared with 2 of 24 reference-first episodes. These counts pool the three models over eight starts

Table 8: Exploratory D–I gaps in percentage points, with 95% intervals obtained by resampling paired physical starts. Positive values favor the direct instruction. No new multiplicity-adjusted significance claims are made.

Model	At any time	At episode end
All models	16.2 [11.7, 20.4]	14.4 [10.0, 18.9]
Nano	14.1 [7.0, 20.8]	16.1 [4.2, 27.6]
Edge	13.5 [7.8, 19.5]	8.6 [0.0, 17.7]
FLUX	21.1 [10.4, 30.7]	18.5 [9.1, 27.6]




Figure 5: The short direct instruction also outperforms reference-first descriptions on average. Intervals are descriptive paired-start bootstrap intervals. D–I was added after the original analyses.

each. The corresponding target-first/reference-first counts are 8/0 for Nano, 4/0 for Edge, and 8/2 for FLUX. By contrast, cube-behind has a negative S–I gap, and the cube scene has no net S–I gap after averaging its four relations. This pattern motivates testing equivalent descriptions within each relation rather than expecting one prompt form to be uniformly preferable. Per-relation estimates involve only eight physical starts and should be read as descriptive localization of the observed effect.

C.3 Reaching an arrangement and keeping it

Across the 864 episodes, 383 reach a stable placement and 191 also satisfy the episode-end criterion. Thus 192 of the 383 any-time passes (50.1%) do not retain the scored arrangement through the final second. Among the 648 episodes that must establish a new arrangement, 167 reach it and 94 retain it at the end. The remaining 73 account for 43.7% of that group’s any-time passes. For initially-satisfied goals, 97 of 216 episodes preserve a final stable outcome. These are raw counts. Figure 7 uses the scene-balanced rates used in the main paper, so its percentages need not equal the raw fractions.

C.4 Complete per-goal counts

Each table entry gives the number of any-time passes followed by the number of episode-end passes, both out of eight starts. A dagger marks a goal satisfied at all eight initial states. No missing observations are represented as zero. These tables report the original physical outcome predicates, not a newly corrected measure of instruction success.




Figure 6: Per-relation S-I differences pooled over models, with descriptive 95% intervals from the eight physical starts for that scene. Positive values favor target-first wording. Initially-satisfied goals contribute zero to this any-time contrast.




Figure 7: Any-time and episode-end placement rates, averaging goals within each scene and weighting scenes equally. The right panel excludes goals already satisfied at reset. Reaching the scored arrangement and keeping it under continued control are distinct requirements.

Table 9: Nano placement counts. Entries are any-time/end, each out of eight.

Goal	Direct	Target-first	Reference-first
Cube left of bowl	8/2	6/3	0/0
Cube right of bowl [†]	8/3	8/2	8/0
Cube in front of bowl	1/1	1/1	0/0
Cube behind bowl	0/0	0/0	7/6
Butter box on raisin box	6/5	6/5	2/1
Butter box left of raisin box	0/0	0/0	0/0
Butter box right of raisin box [†]	8/5	8/4	8/6
Mustard bottle on raisin box	5/5	6/5	4/2
Mustard bottle left of raisin box [†]	8/5	8/5	8/0
Mustard bottle right of raisin box	7/3	8/6	0/0
Left bowl onto right bowl	0/0	0/0	0/0
Right bowl onto left bowl	1/1	1/1	1/0

Table 10: Edge placement counts. Entries are any-time/end, each out of eight.

Goal	Direct	Target-first	Reference-first
Cube left of bowl	3/1	2/1	0/0
Cube right of bowl [†]	8/4	8/3	8/0
Cube in front of bowl	0/0	1/0	0/0
Cube behind bowl	1/0	0/0	0/0
Butter box on raisin box	3/2	0/0	0/0
Butter box left of raisin box	0/0	0/0	0/0
Butter box right of raisin box [†]	8/5	8/3	8/7
Mustard bottle on raisin box	0/0	0/0	0/0
Mustard bottle left of raisin box [†]	8/3	8/4	8/3
Mustard bottle right of raisin box	4/3	4/4	0/0
Left bowl onto right bowl	0/0	0/0	0/0
Right bowl onto left bowl	2/1	1/1	0/0

Table 11: FLUX placement counts. Entries are any-time/end, each out of eight.

Goal	Direct	Target-first	Reference-first
Cube left of bowl	1/0	0/0	4/3
Cube right of bowl [†]	8/2	8/2	8/8
Cube in front of bowl	0/0	0/0	0/0
Cube behind bowl	2/0	1/0	0/0
Butter box on raisin box	2/2	0/0	0/0
Butter box left of raisin box	0/0	0/0	0/0
Butter box right of raisin box [†]	8/4	8/6	8/3
Mustard bottle on raisin box	7/2	5/2	6/2
Mustard bottle left of raisin box [†]	8/6	8/4	8/0
Mustard bottle right of raisin box	8/8	8/3	2/0
Left bowl onto right bowl	7/6	8/5	1/1
Right bowl onto left bowl	6/0	4/0	4/0

D Execution examples and image context

The main paper keeps the mustard example because it directly illustrates the wording comparison. Here we retain the second example and provide the wider camera context. These examples were selected from the recorded outcome differences before inspecting their videos. They illustrate possible behaviors and are not used to estimate their prevalence.

D.1 An arrangement can be reached by moving the reference object

Both Nano instructions ask the robot to move the cube behind the bowl. In the reference-first rollout, however, the robot holds the bowl while the cube remains on the tabletop. Moving the bowl changes the relative positions enough for the cube-behind check to pass. The target-first rollout does not pass that check. Thus, the reference-first score advantage in this example is not evidence of better instruction following: the predicate credits a relation that became true through manipulation of the wrong object.

D.2 How the execution panels were constructed

The mustard pair uses Edge, the mustard-right goal, and physical start C02. The cube pair uses Nano, the cube-behind goal, and start C01. Each displayed pair shares the same initial physical state and uses a common snapshot time within that pair. Mustard snapshots are at 40 seconds, and cube snapshots are at 25.4 seconds. The original recordings were retrieved and checked against their saved file identities before frame extraction.

Nano · cube behind the bowl




Figure 8: Nano’s cube-behind example, moved from Figure 3 in the original draft. Both outcome snapshots are at 25.4 seconds, when the reference-first run first completes the stable-placement interval. The robot holds the bowl in I while the cube is visible on the table. This arrangement no longer satisfies the episode-end check. The labels describe recorded spatial outcomes, not complete instruction success.

All main-paper and cropped appendix panels use the same image rectangle, taken from the exterior camera’s original 1280×720 frame, followed by a 90-degree rotation for readability. The crop does not change within a pair or depend on the wording outcome. No objects or image content were generated or retouched. Figure 9 shows the corresponding full frames in their native orientation. This preserves scene context and allows readers to see what the close views omit.

Edge: mustard right of the box

Nano: cube behind the bowl





Figure 9: Full exterior-camera frames for the two execution examples. Columns correspond to Edge mustard-right and Nano cube-behind. Rows show the shared initial scene, S, and I. The camera image is shown in its native orientation, so image left/right should not be interpreted as the robot-frame relation. These are executed observations, not model-generated predictions.



Automated labels. Agreement does not establish accuracy.

Positive values favor forecasts. The 95% interval is for the pooled estimate.

Figure 10: Three stages of the forecast diagnostic. (a) The relation assigned by the automated readout. Numbers are counts of the 864 windows. Most labels identify neither the requested relation nor an alternative registered relation. (b) Independent VLM agreement on object identity, direction, and final relation. Agreement does not measure accuracy. (c) Change in held-out Brier error when forecast features are added. Positive values favor forecasts. Only the pooled estimate has a conditional 95% interval. Complete export coverage does not establish reliable interpretation or added diagnostic value.

E Can predicted futures help explain an execution failure?

A generated future could make a world–action model easier to diagnose: it might show the wrong object moving before that error appears in the robot’s execution. However, a forecast is useful for this purpose only if its content can be read reliably and it adds information beyond the current scene and proposed actions. We tested this additional diagnostic value on the recorded episodes. The automated readout did not provide a measurable improvement. Low agreement between annotators, together with the absence of independent accuracy checks, leaves the information contained in the videos unresolved.

This was the study’s original primary question. The wording comparisons in the main paper were planned secondary analyses and became the paper’s focus after the behavioral outcomes were available. The original plan specified independent human annotation of generated videos. After viewing those outcomes, we replaced human annotation with two vision–language models (VLMs). This change was recorded before inspecting any VLM outputs. We report the forecast analysis here to preserve the full experimental record, including the limitations of that measurement procedure.

Observable field	Agreement (%)	Cohen’s κ
Moving object	59.1	0.398
Motion direction	48.6	0.108
Visible final relation	20.1	0.077
Reference object for relation	33.3	0.111
Possible release	93.4	0.026
Missing object	97.7	0.000
Hallucinated object	98.5	0.231

Table 12: Independent VLM agreement on the 864 selected forecast windows. These are agreement measures, not accuracy estimates. The original labelers agree on all fields in 5.9% of packets. Agreement is lower on the 862 additional first-request windows: κ is at most 0.114 for every field.

E.1 A forecast and its corresponding action segment

The run records contain 13,968 saved futures across the 864 evaluation episodes. The diagnostic analysis uses one window per episode, selected without consulting the generated video or the later outcome. We select the first request for which the gripper is within 10 cm of the surface of a movable task object. If no such request occurs, we use the episode’s first request. This fallback was needed in two Nano episodes.

Each policy executes a chunk of up to 32 actions at 15 Hz. The exported prediction comes from the same request as that action chunk. The annotation pipeline retains frames assigned to the executed part of the chunk, avoiding a comparison between an unexecuted suffix and behavior after a later replan. The implementation assigns forecast frame k to $k/15$ seconds. Complete checkpoint-specific documentation establishing this temporal and camera correspondence remains to be attached. Therefore, the saved-window comparisons should not be interpreted as a fully validated frame-by-frame forecast accuracy benchmark.

All 864 selected forecasts decoded. We also prepared the first-request clip when it differed from the selected window, yielding 1,726 annotation packets in total. A packet contains an initial image with numbered objects and camera-orientation arrows, followed by the generated sequence. The VLM sees eight evenly spaced frames, including the last retained frame. Frames are presented in a common layout: a wrist-camera view above two exterior views. The instruction, model name, requested goal, wording condition, execution video, and task outcome are hidden. A common layout removes a simple model-identity cue, although visual rendering style may still reveal the generating model.

E.2 What the automated annotators could read

Two independent labelers, Qwen3-VL-235B-A22B-Instruct-FP8 and GLM-4.5V-FP8, identify the moving object, motion direction, visible final relation, reference object, possible release, and missing or hallucinated objects. The object labels are defined in the initial scene. Intended-goal compatibility is computed only after annotation. Unknown and ambiguous observations remain explicit labels. A video cannot establish contact forces or hidden three-dimensional containment from RGB alone.

When the labelers disagree, Qwen adjudicates between their two values and an unknown option, with the candidate order randomized. This procedure produces a complete derived table, but it does not turn agreement into ground truth: the adjudicator is also one of the original labelers. Both original label sets and raw responses are retained.

Agreement is weakest for the quantities needed to infer a destination (Table 12). Only 5.9% of selected-window packets agree on every field. High raw agreement on possible release is also misleading because releases are rare. Its chance-corrected agreement is near zero. Neither averaging the annotators nor retaining only their agreements would by itself establish annotation accuracy.

The adjudicated labels describe upward-only motion in 673 of 864 windows (77.9%). Only 39 windows (4.5%) are assigned the requested final relation, 99 (11.5%) another registered relation, 608 (70.4%) neither, and 118 (13.7%) an unknown relation. These are properties of the automated

Comparison	Events / windows	Baseline	+ Forecast	ΔB
All models	47 / 864	0.04012	0.04085	−0.00073
Cosmos3 Nano	15 / 288	0.04012	0.04151	−0.00139
Cosmos3 Edge	8 / 288	0.02491	0.02530	−0.00040
FLUX 3 Action	24 / 288	0.06210	0.06104	+0.00106

Table 13: Held-out event prediction using the automatically labeled forecasts. Lower Brier error is better. Positive ΔB favors including forecast labels. The pooled 95% interval is $[-0.00177, +0.00016]$, obtained by resampling initial states within scenes while keeping the fitted held-out losses fixed. Model-specific estimates are descriptive. No model-specific intervals were produced in the saved analysis.

readout, not independently established properties of the generated videos. They suggest a practical obstacle: a short clip dominated by grasping and lifting may reveal object selection while providing little visible evidence of the eventual placement. That interpretation needs confirmation from reliable clip annotations.

E.3 Testing information beyond the current state and actions

The target is a detected wrong-object manipulation or a contradictory released placement in the selected action chunk. The event scorer marks 47 of 864 chunks positive: 14 wrong-object detections and 33 contradictory-release detections. The release rule can trigger on a contact-to-detached transition without requiring substantive prior manipulation. These detections have not all been checked against raw recordings, so they should not be equated with 47 independently verified semantic errors.

The baseline is an L2-regularized logistic classifier. Its inputs describe the model, scene, goal, current object positions, distance to the goal region, whether the goal already holds, and gripper state. They also include distance and nearest-object identity along forward kinematics of the proposed actions, together with prior contact and lift indicators. This is a privileged simulator diagnostic, not a deployable robot failure detector. It excludes wording form, actual future state, and final task success. The augmented classifier adds only the blinded forecast labels, including explicit unknown categories.

We hold out one initial-state slot at a time across all scenes, models, goals, and wording forms, giving eight folds. Feature standardization and categorical encoding are fitted on the training fold. Both classifiers use fixed regularization ($C = 1$), with no hyperparameter search. The score is the difference in held-out Brier error,

$$\Delta B = B(\text{state and actions}) - B(\text{state, actions, and forecast labels}),$$

so a positive value indicates added diagnostic information. Models receive equal weight, as do scenes within each model and goals within each scene.

Adding forecast labels does not improve the pooled held-out score (Table 13). The small positive FLUX estimate is insufficient to establish a model-specific benefit. The interval is conditional on the fitted diagnostics and the chosen labels. It does not capture annotation error or variability from refitting the diagnostic in each resample.

A missingness-only control changes Brier error by +0.000011. The implemented no-motion/current-goal control changes it by +0.000009: it supplies synthetic no-motion labels and the privileged current-goal flag, rather than labeling a repeated current image through the VLM pipeline. We describe it accordingly instead of treating it as the originally proposed video-persistence baseline. Shuffling forecast associations across states over 100 seeds gives a mean change of −0.00053, with a 2.5th–97.5th percentile range of $[-0.00177, +0.00049]$. This is the permutation distribution, not a second confidence interval for the observed effect, and the saved analysis did not calculate a formal permutation-test p -value.

E.4 What this result tells us about useful predictions

The current procedure provides no demonstrated incremental diagnostic benefit from the automatically extracted forecast features. It does not establish that generated futures are uninformative. Two explanations remain compatible with the evidence: the short forecast adds little beyond state and action, or the annotation pipeline fails to recover useful information that is present in the video. Unverified temporal correspondence and event detections introduce further uncertainty. Complete video export is therefore insufficient evidence of diagnostic usefulness.

There is a suggestive descriptive signal in object selection. Windows labeled as moving the reference object are followed by a detected event in 23 of 247 cases (9.3%), compared with 19 of 540 (3.5%) when the requested object is labeled as moving. This comparison was examined after observing the results. It is unadjusted for scene, goal, and model, uses noisy labels, and does not establish incremental information over the baseline. It motivates checking object identity more carefully, not a claim that forecasts predict failure reliably.

The next step is to validate measurements on existing recordings. Blinded human annotations or verified tracking can establish accuracy for object identity, motion, and destination. Agreement between VLMs is an uncertainty signal, not ground truth. Clips need verified timing and camera mappings, and release detections need checks for contact-reset artifacts. Testing windows that expose an object or destination choice, with the selection rule fixed beforehand, can then distinguish an uninformative forecast from an inadequate readout.