

Cosmos3 Edge: forecast and execution of the same 32-action chunk


Episode RWS-E3-S1-C01-B-D · annotation packet Aef9bebaab690 · request 1 (chunk 2) · seed 6100

Instruction: “Put the rubiks cube behind the bowl”

Request input: tick 32, t = 2.133 s · 32 actions executed at ticks 32–63 · after the last: tick 64, t = 4.267 s

Exterior camera over_shoulder_left_camera = rows 360–539, columns 0–319 of the 640×540 policy input

Initial observation: request-1 input, tick 32



Original VLM labels of this forecast

1 = banana 2 = bowl (reference)

3 = Rubik's cube (mover)

	Moving object	Direction	Final relation†	Relation object
A (Qwen3-VL)	3 (cube)	up only	none	none
B (GLM-4.5V)	3 (cube)	up only	on top of	none
Adjudicated	3 (cube)	up only	none	none
<i>Derived</i>	<i>mover</i>	<i>up only</i>	<i>not goal</i>	–

† A and B disagreed on this field; the adjudicator (Qwen3-VL) chose the value shown.

k = 0 · +0.00 s

tick 32 · t = 2.133 s

k = 8 · +0.53 s

tick 40 · t = 2.667 s

k = 16 · +1.07 s

tick 48 · t = 3.200 s






k = 24 · +1.60 s

tick 56 · t = 3.733 s




k = 32 · +2.13 s

tick 64 · t = 4.267 s

PREDICTED generated frame k of the request-1 forecast



EXECUTED simulator frame at tick 32 + k, after k of the 32 executed actions



All ten panels show the same 160×90-pixel window (dashed box) of the 320×180 exterior view in the composite given to the policy; source pixels are not resampled. Column k pairs generated frame k with the executed frame after k actions (tick 32 + k), the nominal 15 Hz mapping defined by the model and export code. Generated frame 0 is the model's reconstruction of its input. The mapping does not show that the predicted motion keeps the executed pace.

Labelers A (Qwen3-VL-235B-A22B-Instruct-FP8) and B (GLM-4.5V-FP8) saw only legend.png and generated frames k = 0, 5, 9, 14, 18, 23, 27, 32, never the instruction or the execution; disagreements went to an adjudicator (Qwen3-VL-235B-A22B-Instruct-FP8) with the same images. A, B and the adjudication all answer no / no / no for possible release / missing object / hallucinated object. Derived: the labels mapped to scene roles after unblinding.

Example chosen by the manifest rule “Lexicographically first primary-window episode for model; selected before media inspection” and kept regardless of what it shows.

Cosmos3 Edge decodes rows 0–527 of the 540-row composite; the missing rows lie outside this window.