

Cosmos3 Edge: original annotation packet Aef9bebaab690


Episode RWS-E3-S1-C01-B-D · request 1, input at tick 32 (t = 2.133 s)
33 generated frames, nominally 15 Hz · prediction only
Instruction (hidden from the labelers): “Put the rubiks cube behind the bowl”

legend.png, as given to the VLM labelers (shown at 1/2 resolution)



Episode's first observation (tick 0, before any action). 1 = banana 2 = bowl (reference) 3 = Rubik's cube (mover)
Yellow circles and numbers mark the objects; each number is drawn above and to the right of its circle, and the direction arrows are drawn over the numbers. In the left view the arrow labels cover 1 and 3, and the bowl's 2 falls beside the banana; in the right view the cube's 3 falls on the banana's circle. Both examples have byte-identical legends (same restored initial state).

contact_sheet.png, the original human-review summary (not shown to the labelers)



Tiles: generated frames $k = 0, 4, \dots, 32$ of the request-1 forecast, every second pixel; in each, the wrist view is above the left | right exterior views.

+s is the nominal time after the request input: $t = 2.133 \text{ s} + k/15$.

The labelers instead saw legend.png and frames $k = 0, 5, 9, 14, 18, 23, 27, 32$ of forecast.mkv at 640x544.

Prediction only: nothing here is paired with the execution. Frames 0, 8, 16, 24 and 32 appear with the matching executed frames in the companion figure, cut from the original future array.

The packet layout stretches this model's 168 decoded exterior rows to 184 (one layout for all models).